


~~Monty Python's~~ Stephan's search for the holy grail

1. Starting point: Fred's algorithm
2. Improvements, where and why
3. Hierarchical clustering → Variational Bayes
4. PMC → Variational Bayes


Fred's algorithm



Overall goal: Find “good” proposal for Importance Sampling


$$\int P(x) = \int \frac{P(x)}{q(x)} q(x) \approx \sum_{x_i \in \text{samples}} \frac{P(x_i)}{q(x_i)} \quad \text{where } x_i \sim q(x)$$

Fred's algorithm



does its job


Fred's algorithm



does its job

does not reduce
redundant
information

Fred's algorithm



does its job

does not reduce
redundant
information

no re-use of samples
in subsequent
update steps

<https://github.com/fredRos/pypmc>

```
'''This example illustrates how to run a Markov Chain using pypmc'''
```

```
import numpy as np
import pypmc
```

```
# define a proposal
```

```
prop_dof = 50.
prop_sigma = np.array([[0.1, 0. ],
                        [0. , 0.02]])
prop = pypmc.density.student_t.LocalStud
```

```
# define the target; i.e., the function
# In this case, it is a Gaussian with me
# covariance "target_sigma".
```

```
# Note that the target function "log_target"
# unnormalized gaussian density.
```

```
target_sigma = np.array([[0.01, 0.003 ],
                          [0.003, 0.0025]])
```

```
inv_target_sigma = np.linalg.inv(target_sigma)
target_mean = np.array([4.3, 1.1])
```

```
def unnormalized_log_pdf_gauss(x, mu, inv_sigma):
```

```
    diff = x - mu
    return -0.5 * diff.dot(inv_sigma).dot(diff)
```

```
log_target = lambda x: unnormalized_log_pdf_gauss(x, target_mean, i
```

```
# choose a bad initialization
```

```
start = np.array([-2., 10.])
```

```
# define the markov chain object
```


```
mc = pypmc.sampler.markov_chain.AdaptiveMarkovChain(log_target, pro
```

```
# run burn-in
```

```
mc.run(10**4)
```

```
# delete burn-in from history
```

```
mc.history.clear()
```



(effective samples).

- **rel_tol** -

Relative tolerance ϵ . If two consecutive values of the log likelihood bound, L_t, L_{t-1} , are close, declare convergence. More precisely, check that

$$\left\| \frac{L_t - L_{t-1}}{L_t} \right\| < \epsilon.$$

- **abs_tol** -

Absolute tolerance ϵ_a . If the current bound L_t is close to zero, ($L_t < \epsilon_a$), declare convergence if

$$\|L_t - L_{t-1}\| < \epsilon_a.$$

- **verbose** - Output status information after each update.

set_variational_parameters()

Reset the parameters to the submitted values or default.

Use this function to set the prior value (indicated by the subscript θ as in α_θ) or the initial value (e.g., α) used in the iterative procedure to find the posterior value of the variational distribution.

Every parameter can be set in two ways:

1. It is specified for only one component, then it is copied to all other components.
2. It is specified separately for each component as a K vector.

The prior and posterior variational distributions of μ and Λ for each component are given by

$$q(\mu, \Lambda) = q(\mu|\Lambda)q(\Lambda) = \prod_{k=1}^K \mathcal{N}(\mu_k | \mathbf{m}_k, (\beta_k \Lambda_k)^{-1}) \mathcal{W}(\Lambda_k | \mathbf{W}_k, \nu_k),$$

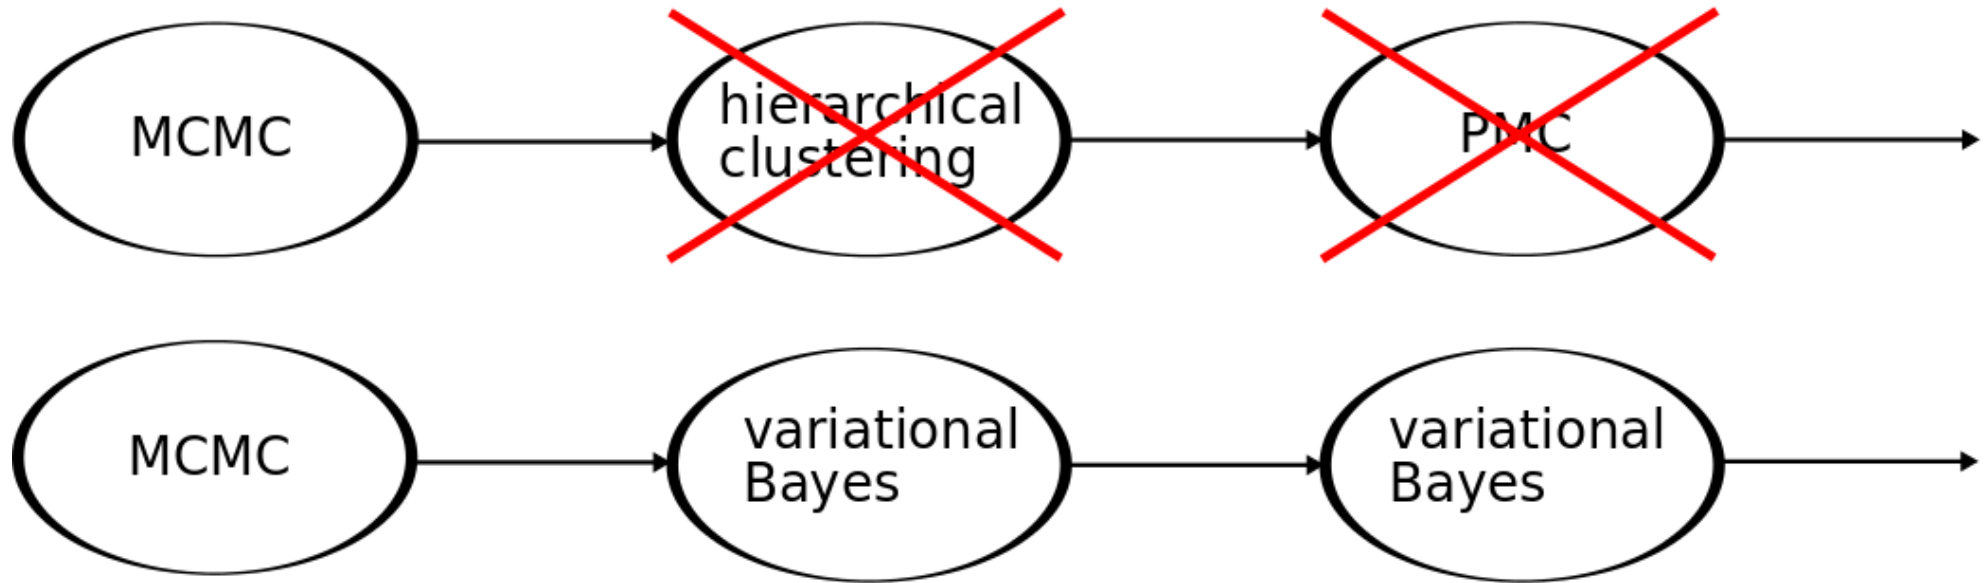
where \mathcal{N} denotes a Gaussian and \mathcal{W} a Wishart distribution. The weights π follow a Dirichlet distribution

$$q(\pi) = \text{Dir}(\pi | \alpha).$$

Warning

This function may delete results obtained by `update()`.

Suggested new algorithm



Variational Bayes

- adapt Gaussian mixture to samples
- Expectation Maximization (EM) with non trivial prior
- Very detailed derivation in “Pattern Recognition and Machine learning” (Christopher M. Bishop)

Hierarchical clustering (hc) → VB

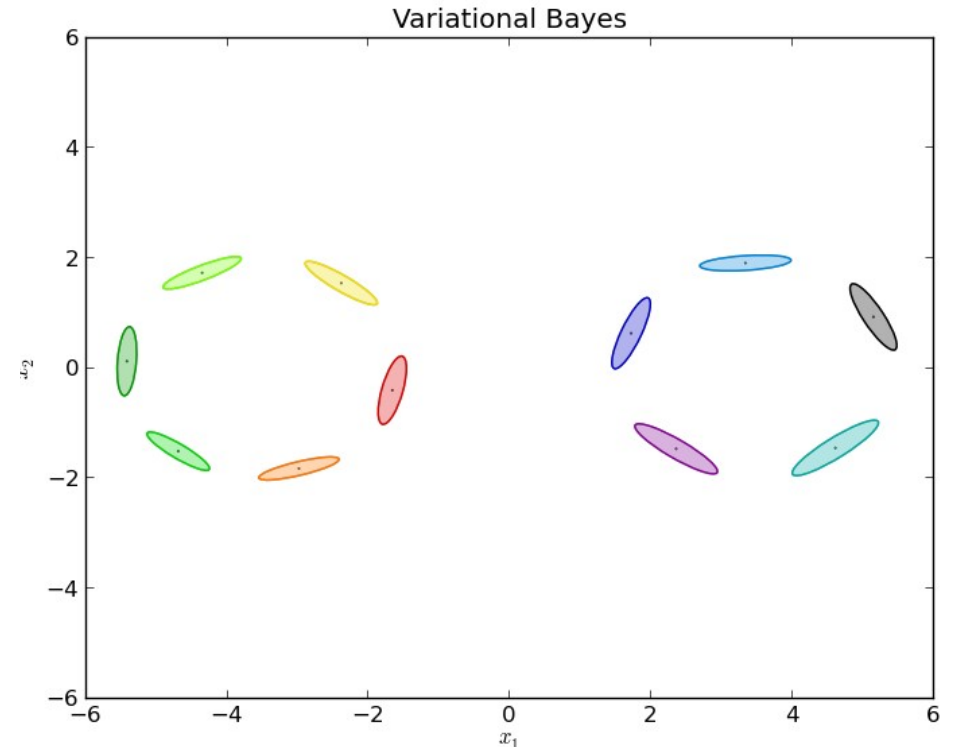
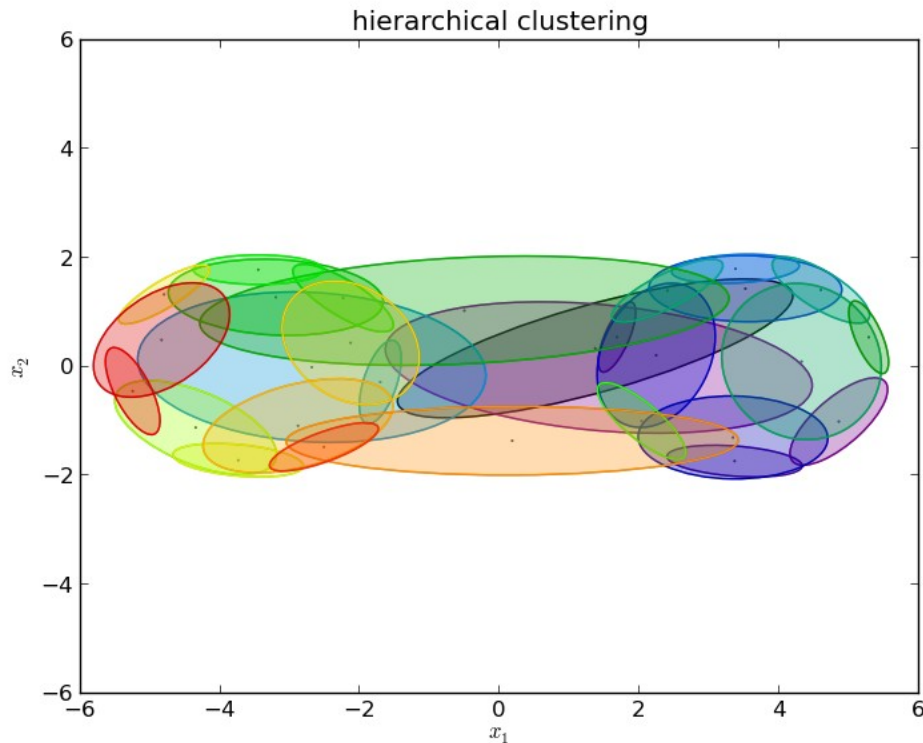
Advantages:

- Can at least reproduce the hc results, most runs bring better perplexity and effective sample size
- Automatically removes components with redundant information

Disadvantage:

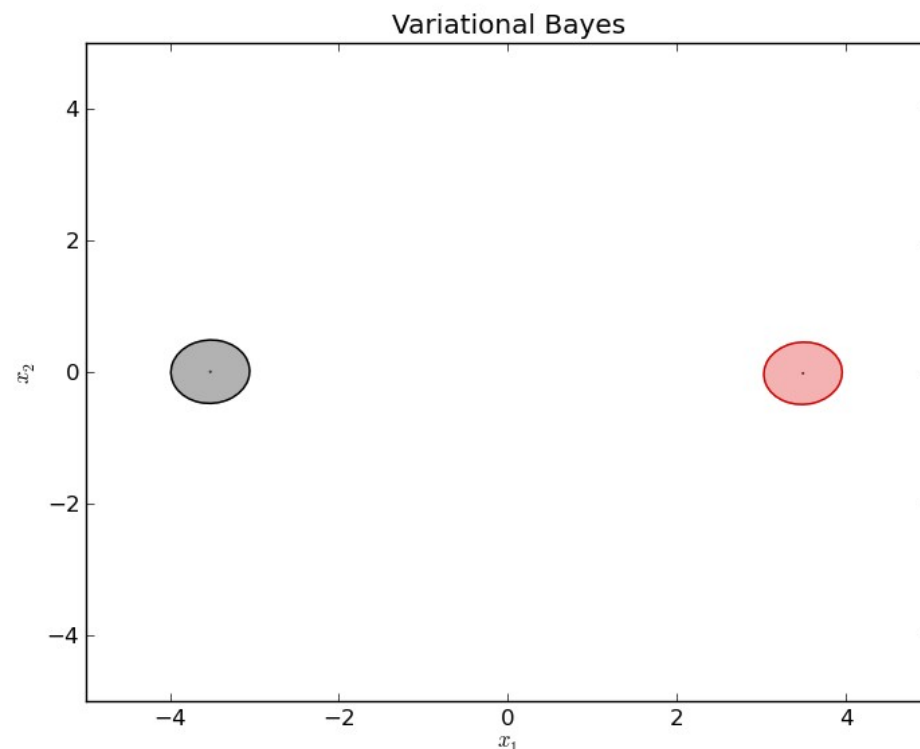
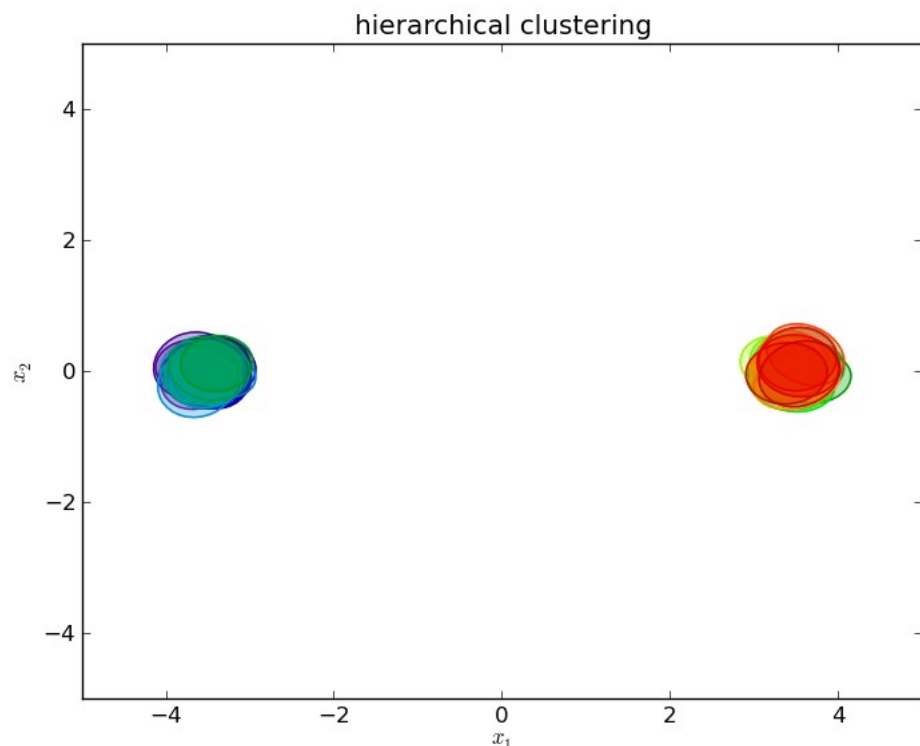
- Computationally much more expensive

Example: Gaussian shells, 2 dim



- Start clustering with 15 components per shell
 - Converged with 29 components after 21 steps
 - Perplexity: ~32%
 - ESS: ~27%
- Converged with 11 components after 304 steps
 - Perplexity: ~65%
 - ESS: ~46%

Example: Gaussian shells, 20 dim



- Start clustering with 25 components per shell
 - Converged with 50 components after 5 steps
 - Perplexity: ~33%
 - ESS: ~23%
- Converged with 2 components after 304 steps
 - Perplexity: ~39%
 - ESS: ~28%

PMC \rightarrow VB (to come ...)


Hopes:

- VB allows (forces) you to define priors for the model parameters
 - \rightarrow We can set the posterior from the previous step as new prior. \rightarrow Use all information properly.

Fears:

- Unjustifiable computational effort

Timeline



Variational Bayes

Given data \mathbf{X} , we can write for the evidence of any model:

$$\ln p(\mathbf{X}) = \mathcal{L}(q) + KL(q||p)$$

where we define:

$$\mathcal{L}(q) \equiv \int q(\mathbf{Z}, \boldsymbol{\theta}) \ln \left\{ \frac{p(\mathbf{X}, \mathbf{Z}, \boldsymbol{\theta})}{q(\mathbf{Z}, \boldsymbol{\theta})} \right\} d\mathbf{Z} d\boldsymbol{\theta}$$

$$KL(q||p) \equiv - \int q(\mathbf{Z}, \boldsymbol{\theta}) \ln \left\{ \frac{p(\mathbf{Z}, \boldsymbol{\theta} | \mathbf{X})}{q(\mathbf{Z}, \boldsymbol{\theta})} \right\} d\mathbf{Z} d\boldsymbol{\theta}$$

$\boldsymbol{\theta}$ are model parameters

\mathbf{Z} are “latent variables” (next slide)

Variational Bayes

Definition: Latent (hidden) variables

Be $\mathbf{X} = \{x_1, \dots, x_N\}$ your data

$\mathbf{Z} = \{z_1, \dots, z_N\}$ is called latent if \mathbf{Z} contains information which can not uniquely be determined from \mathbf{X} but inferred in a probabilistic model.

Example (Gaussian mixtures):

$$g(x) = \sum_i \omega_i \mathcal{N}(x | \mu_i, \Sigma_i)$$
$$\mathbf{X} \sim g(x)$$

Then the index i is a latent variable

Variational Bayes

Given data \mathbf{X} , we can write for the evidence of any model:

$$\ln p(\mathbf{X}) = \mathcal{L}(q) + KL(q||p)$$

Holds for any probability distribution q

Main idea: Restrict q to factorize as below and minimize the Kullback-Leibler divergence

$$q(\mathbf{Z}, \boldsymbol{\theta}) = q(\mathbf{Z})q(\boldsymbol{\theta})$$

I.e. we look for the closest factorizing q to p in the sense of $KL(q||p)$

Variational Bayes

Given data \mathbf{X} , we can write for the evidence of any model:

$$\ln p(\mathbf{X}) = \mathcal{L}(q) + KL(q \| p)$$

Minimizing KL is equivalent to maximizing \mathcal{L} .

Maximizing \mathcal{L} , in the case of a Gaussian mixture model, leads to an EM-like algorithm.